





# L7: Multilevel (Mixed) Models




# Grouped data



# Grouped data



Data = Mean + Treatment effect + Error



Data = Mean (+ Treatment effect) + Error  
(+ Group treatment effect) + Group error

# Up to now...

- We have assumed that factors in our GLMs have had their levels chosen, or “fixed”
  - we assumed that all error lies in our dependent variable
  - we deliberately choose Drugs A, B and C to compare – we don’t pick three drugs at random
- So our results apply **only** to the exact levels we chose - Drugs A vs. B vs. C – no inference about “Drugs” in general.



But suppose we want to know about effects of a random sample of treatments...

## Fixed Effects

- the model estimates the mean effect of each factor level on the response
- for treatments you specified in your hypothesis
- coefficients will be means for each group

"no pooling"

"partial pooling"

## Random Effects

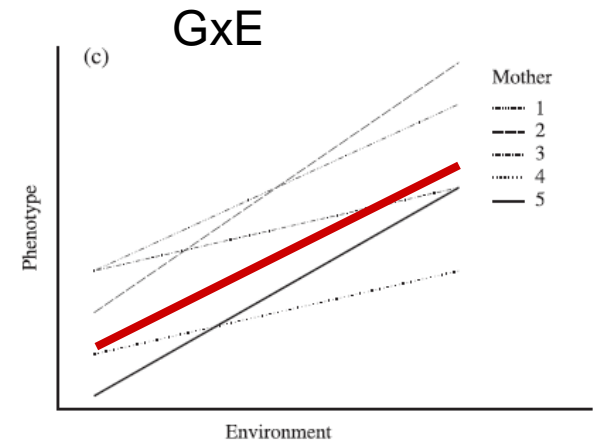
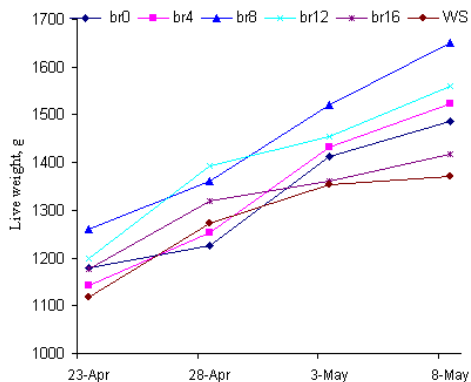
- the model can also estimate **variance** of responses among the factor levels
- for treatments you didn't specify or control
  - Species, Genotype, Country
- for groupings in your data
  - Location, Individual, Brood
  - Block, Field, Plot, Sub-plot
- estimates for each group  $\neq$  means; drawn towards grand mean according to within-group variability

# More formally...

- A model with **a single fixed** factor estimates:
  - $y_{ij} = \mu + \alpha_i + \varepsilon_{ij}$     e.g. height = mean + sex + error
  - $\mu$  = mean,  $\alpha_i$  = treatment effect,  $\varepsilon_{ij}$  = residual ( $N(0, \sigma^2)$ )
- A model with **a single random factor** predicts the variance from which each factor level is drawn:
  - $y_{ij} = \mu + A_i + \varepsilon_{ij}$     e.g. growth rate = mean + individual + error
  - $A$  is the random effect and is assumed  $N(0, \sigma^2)$ ; each individual  $i$  is drawn at random from this distribution
  - There is *sampling error* associated with the choice of individual ( $A$ ) and with measuring growth rate, and the errors are independent
- A **mixed model** combines fixed and random effects, e.g:
  - $y_{ij} = \mu + \alpha_i + A_i + \varepsilon_{ij}$     e.g. growth rate = mean + sex + individual + error

# Two uses of random factors

1. Where the variance of a factor's effects is more interesting than knowing the actual means.
  - e.g. Nussey *et al* (2005) J. Anim. Ecol. estimating among-individual variation in plasticity (response of calving date to climate)
2. Where you want to know mean effects of a factor's levels but want them to fit a distribution (central tendency)
  - e.g. known genotypes, species, other biological groups.



# Examples of groups treated as adding random variation:

- Groups of individuals (additional error associated with group)
  - Broods/families
  - Populations
- Repeated measurements on individuals (additional error associated with choice of individual)
  - E.g. growth rate
- Grouping in time or space (additional error associated with choice of time or place)
  - Day of measurement/climate
  - Block (field/plot/greenhouse/incubator/waterbath/experimenter)
  - Site, Region, County, Country, etc
- Populations sampled at random, where you want to make inferences about the variance among populations as well as an overall mean



# “Mixed Effects”: Fixed + Random

## Douglas Bates

If levels of a factor represent a **random sample** from the set of all possible levels, we model them using **random effects**. Conclusions apply to the population of levels from which we chose. Sampling at random means there will be **sampling error** associated with a predictor variable, which violates the GLM assumption that  $X$  is measured without error.



“ *Mixed-effects models* or, more simply, *mixed models* are statistical models that incorporate both fixed-effects parameters and random effects. Because of the way that we will define random effects, a model with random effects always includes at least one fixed-effects parameter. Thus, any model with random effects is a mixed model. ”

There are two things to notice about this distinction between fixed-effects parameters and random effects. First, the names are misleading because the distinction between fixed and random is more a property of the levels of the categorical covariate than a property of the effects associated with them. Secondly, we distinguish between “fixed-effects parameters”, which are indeed parameters in the statistical model, and “random effects”, which, strictly speaking, are not parameters. As we will see shortly, random effects are unobserved random variables.

# What is a “Random Effect”?

Andrew Gelman (2005: *Ann. Stat.* 33:1–53)



5 definitions of “random effects” (mutually inconsistent):



1. Fixed effects are constant across individuals, and random effects vary. For example, in a growth study, a model with random intercepts  $a_i$  and fixed slope  $b$  corresponds to parallel lines for different individuals  $i$  ( $y_{it} = a_i + b_t$ ). Kreft and De Leeuw (1998).
2. Effects are fixed if they are interesting in themselves or random if there is interest in the underlying population. Searle et al (1992).
3. "When a sample exhausts the population, the corresponding variable is *fixed*; when the sample is a small part of the population the corresponding variable is *random*." (Green and Tukey, 1960)
4. "If an effect is assumed to be a realized value of a random variable, it is called a random effect." (LaMotte, 1983)
5. Fixed effects are estimated using least squares (or, more generally, maximum likelihood) and random effects are estimated with shrinkage ("linear unbiased prediction" in the terminology of Robinson, 1991). This definition is standard in the multilevel modeling literature (see e.g. Snijders and Bosker, 1999, Section 4.2)

→ Just refer to levels of grouping in multi-level / hierarchical models

# Hierarchical grouping ("nesting", stratified sampling)

You may have replicated samples within units of space or time: grouped replicates likely to be more similar

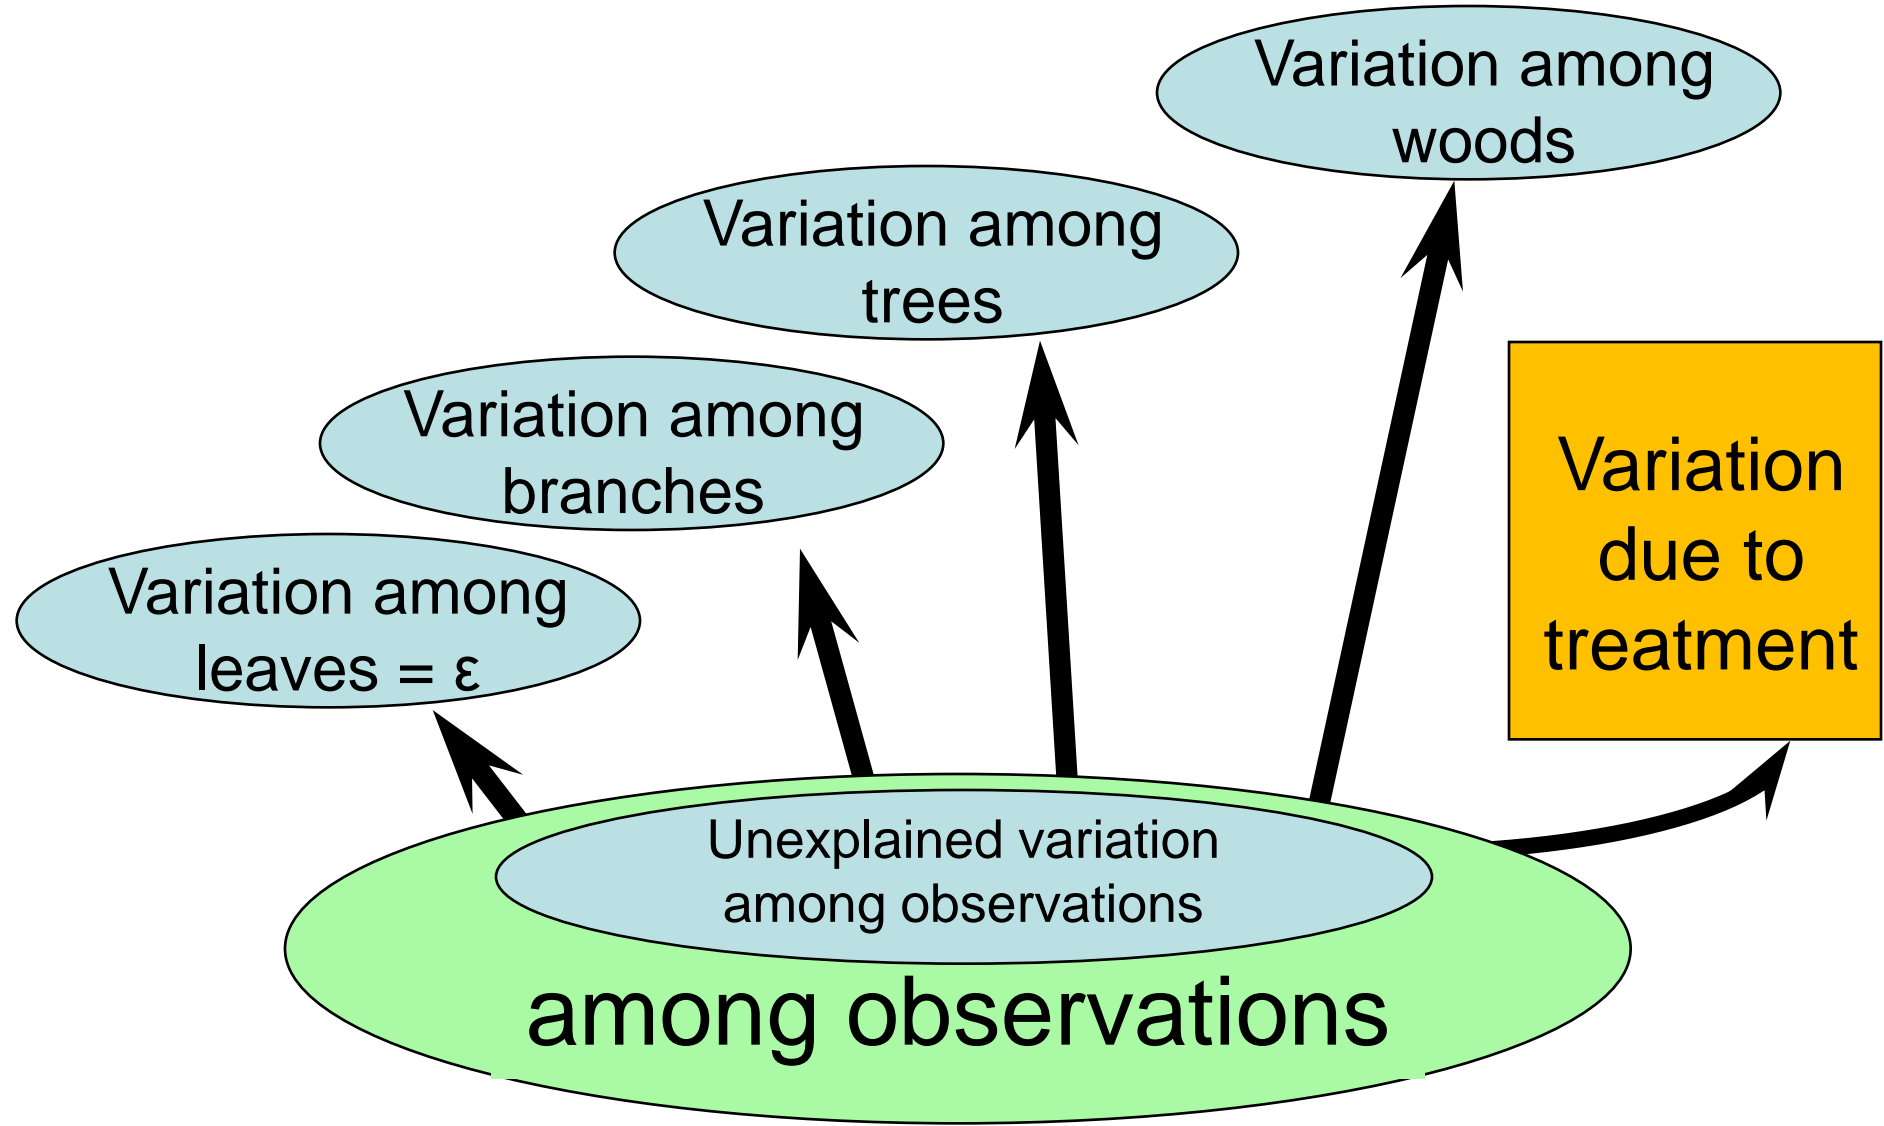
- Chicks in a nest
- Individuals within families
- Sizes at ages (repeated measures)
- Samples within plots within fields within farms within landscapes within regions...
- Locations within tables within glasshouses



*“**Closer** replicates being **more** similar” will crop up for time series and spatial models in the concept of **autocorrelation***

# An example of multiple error terms

- Observations = insect leaf damage on trees



# Example 1: Great tit nest boxes




Experiment to look at influence of heating or cooling on weights of chicks

- 12 nest boxes randomly assigned to 3 treatments (control, warm, cool)
- Fixed effects are treatment (of box) and sex (of individual)
- Random effect (grouping factor) is nest box
- Gives two “strata” for errors:
  - Nest box (group) and chicks within nest (individual)

# Ignoring nest box

“complete pooling”

```
xyplot(WEIGHT ~ SEX | TREAT, auto.key= TRUE)
```



# Ignoring nest box

“complete pooling”

```
> regstyle <- lm(WEIGHT~SEX*TREAT)
> summary(regstyle)
```

Call:

```
lm(formula = WEIGHT ~ SEX * TREAT)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.09080	-0.28797	0.06284	0.36663	1.01295

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.899443	0.138885	71.278	< 2e-16 ***
SEXMale	0.182225	0.196413	0.928	0.356911
TREATCooled	0.008162	0.196413	0.042	0.966978
TREATWarmed	0.283346	0.196413	1.443	0.153861
SEXMale:TREATCooled	-0.966490	0.277770	-3.479	0.000895 ***
SEXMale:TREATWarmed	-0.067498	0.277770	-0.243	0.808760

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4811 on 66 degrees of freedom

Multiple R-Squared: 0.4067, Adjusted R-squared: 0.3618

F-statistic: 9.05 on 5 and 66 DF, p-value: 1.371e-06

Residual Variance =  $0.48^2 = 0.23$

# Ignoring nest box

“complete pooling”


## Fixed Effects model

```
> anova(regstyle)
```

Analysis of Variance Table

Response: WEIGHT

	Df	Sum Sq	Mean Sq	F value	Pr(>F)						
SEX	1	0.4749	0.4749	2.0519	0.156738						
TREAT	2	6.5053	3.2527	14.0522	8.238e-06 ***						
SEX:TREAT	2	3.4937	1.7468	7.5468	0.001118 **						
Residuals	66	15.2770	0.2315								
---											
Signif. codes:	0	***	0.001	**	0.01	*	0.05	.	0.1		1




Total df recorded as 71, with 66 in error:  
too high, since chicks within each  
nestbox are not independent



# What should we do?

partial pooling

Sex of individual



$$\text{Data} = \text{Mean} + \text{Treatment} + \text{Error} \\ + \text{Group Treatment} + \text{Group error}$$



Temperature of box


Using the correct error structure should give optimal estimates of each treatment's effects (as well as estimating variance among nest boxes).

```
Great.tits <- groupedData(WEIGHT~SEX|BOX, outer=TREAT)
```

```
plot(Great.tits)
```

*NB: can specify grouping structure as a `groupedData()` object before analysis to access complicated graphs simply*

Factor that groups subjects



```
> modell <- lme(WEIGHT~SEX*TREAT, random=~1|BOX, data=gg)
```

```
> summary(modell)
```

```
Linear mixed-effects model fit by REML
```

```
Data: gg
```

```
      AIC      BIC    logLik
89.34314 106.8604 -36.67157
```

```
Random effects:
```

```
Formula: ~1 | BOX
```

```
(Intercept) Residual
```

```
StdDev: 0.3953251 0.3218725
```

```
Fixed effects: WEIGHT ~ SEX * TREAT
```

	Value	Std.Error	DF	t-value	p-value
(Intercept)	9.899448	0.2184124	57	45.32455	0.0000
SEXMale	0.182225	0.1314039	57	1.38676	0.1709
TREATCooled	0.008162	0.3088818	9	0.02643	0.9795
TREATWarmed	0.283346	0.3088818	9	0.91733	0.3829
SEXMale:TREATCooled	-0.966490	0.1858332	57	-5.20085	0.0000
SEXMale:TREATWarmed	-0.067498	0.1858332	57	-0.36322	0.7178

```
Correlation:
```

	(Intr)	SEXMal	TREATC	TREATW	SEXM:TREATC
SEXMale	-0.301				
TREATCooled	-0.707	0.213			
TREATWarmed	-0.707	0.213	0.500		
SEXMale:TREATCooled	0.213	-0.707	-0.301	-0.150	
SEXMale:TREATWarmed	0.213	-0.707	-0.150	-0.301	0.500

```
Standardized Within-Group Residuals:
```

Min	Q1	Med	Q3	Max
-2.0153036	-0.7082614	-0.0123904	0.5418996	1.9656051

```
Number of Observations: 72
```

```
Number of Groups: 12
```

Fit by Restricted Maximum Likelihood  
(default)

### Goodness of fit measures:

When comparing models, the model with the lower AIC and greater log-likelihood is a better fit

Estimate of the standard deviation of the distribution of random effects

Estimates of the standard deviation of the within-group error distribution:

**residual variance =  $0.32^2 = 0.103$**  (cf 0.23)

Coefficients (default “treatment contrasts”)

Useful to look at to check the model was specified correctly

# Importance of correct model specification

> anova (mixed effects model)

$$F = MS_{\text{effect}}/0.103$$

Analysis of Variance Table

	Df	Sum Sq	Mean Sq	Denom	F value	Pr(>F)	
SEX	1	0.475	0.475	66.000	4.5843	0.03596	*
TREAT	2	0.647	0.324	66.000	3.1237	0.05056	.
SEX:TREAT	2	3.494	1.747	66.000	16.8611	1.216e-06	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

> anova (fixed effects ignoring box)

Analysis of Variance Table

Response: WEIGHT

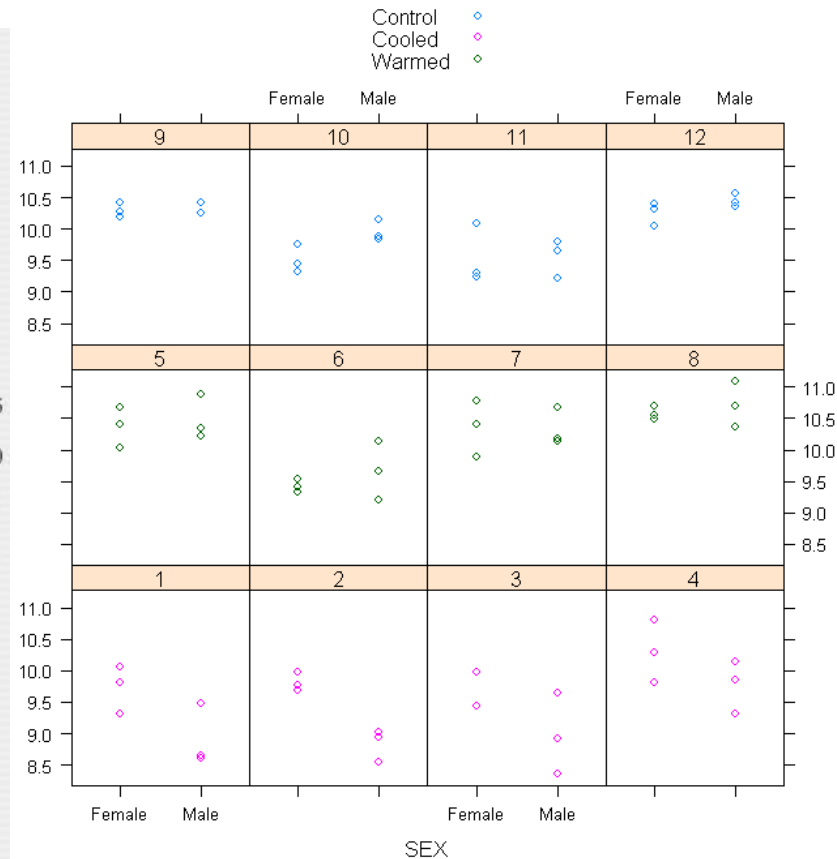
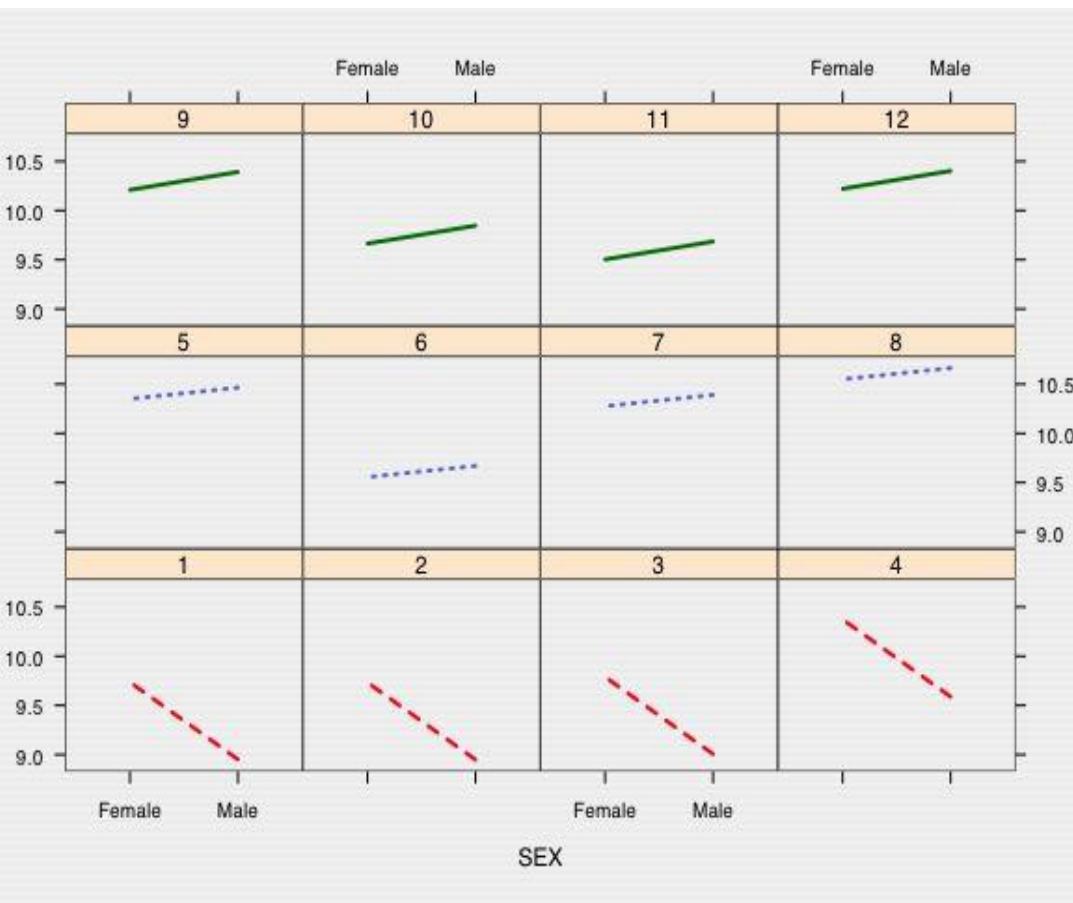
$$F = MS_{\text{effect}}/0.2315$$

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
SEX	1	0.4749	0.4749	2.0519	0.156738	
TREAT	2	6.5053	3.2527	14.0522	8.238e-06	***
SEX:TREAT	2	3.4937	1.7468	7.5468	0.001118	**
Residuals	66	15.2770	0.2315			


---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

# So how well do the predictions fit?




# Checking assumptions




Homoscedasticity

```
plot(modell1)
```



Normality of within-group error

```
qqnorm(modell1)
```



Normality of random effects

```
qqnorm(modell1, ~ranef(.))
```

# Model outputs

Other useful functions:

- **VarCorr(model)** will give you the variance and standard deviations of the distributions of random effects and residuals (N.B. you already get the SDs in the `summary()` output)
- **ranef(model)** will give you estimated coefficients for whatever varies among factor levels. This is useful to check you have specified the model correctly.
- **intervals(model)** will provide 95% C.I. for coefficients of the fixed effects and the standard deviations of the random effects and residual distribution. This is useful to check the precision of the estimates; huge intervals can indicate misspecification of the model.

# Packages in R

- `library(nlme)` for **`lme()`**
  - Older package
  - Allows autocorrelation structures...
- `library(lme4)` for **`lmer()`**
  - Faster, better?
  - Allows 2+ non-nested groupings
  - Allows generalised models

## Fitting linear mixed models

Using the `lme4` package

*by Douglas Bates*

The `lme` function, which fits linear mixed models of the form described in [Pinheiro and Bates \(2000\)](#), has been available in the required R package `nlme` for several years. Recently my colleagues and I have been developing another R package, called `lme4`, and its `lmer` function which provides more flexible fitting of linear mixed models and also provides extensions to generalized linear mixed models.

The good news for `lme` users is that the `lmer` function fits a greater range of models, is more reliable, and is faster than the `lme` function. The bad news is that the model specification has been changed slightly. The purpose of this article is to introduce `lmer`, to describe how it can be used to fit linear mixed models and to highlight some of the differences between `lmer` and `lme`.

---

```
> lme1 <- lme( y~x, random= ~1|z ) is equivalent to
> lmer1 <- lmer( y~x + (1|z) )
```



# Specifying models in **lmer()**

FIXED (treatment) effects

( +  $\epsilon$  )

```
mod031 <- lmer( No ~ R + HC + M + CT + L +  
  (1|Year:S) + (1|C) + (1|YF:FID) + (1|YF:FID:L), data=ins)
```

Survey within year

Field within farm

Cluster (neighbourhood)

Location within field within farm

RANDOM effects (grouping structure)

lmer() formula:

| introduces a grouping variable

/ indicates nesting

: indicates both an interaction and nesting;

a + b%in%a means a + a:b

# Instead of ML, use REML for fitting

- **Maximum likelihood** is the mathematical technique of estimating the parameters that would make the data most likely (*e.g. if you toss a coin 1000 times and get 509 heads, a coin with a 0.509 chance of heads would be most likely to produce that*).
- **Restricted maximum likelihood** “was developed... because maximum likelihood estimates of variance components take no account of the degrees of freedom used in estimating treatment effects, [so] they have a downwards bias which increases with the number of fixed effects in the model. This in turn leads to under-estimates of standard errors for fixed effects.”

<http://www.vsni.co.uk/products/genstat/htmlhelp/server/REML.htm>

# Hypothesis testing 1. Fixed effects

## 1. *F*-tests (recommended by Bates). Conservative

```
# Use anova() to test the significance of fixed terms by an F-test  
> anova(model1)
```

	numDF	denDF	F-value	p-value
(Intercept)	1	57	6797.970	<.0001
SEX	1	57	4.584	0.0366
TREAT	2	9	3.124	0.0933
SEX:TREAT	2	57	16.861	<.0001

## 2. Likelihood-Ratio tests (recommended by Crawley). Anti-conservative (note method=ML)

```
> model2 <- lme( WEIGHT~SEX*TREAT, random=~1|BOX, data=gg, method="ML" )  
> model3 <- lme( WEIGHT~SEX+TREAT, random=~1|BOX, data=gg, method="ML" )  
> anova(model2,model3)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
model2	1	8	78.25097	96.4643	-31.12548			
model3	2	6	102.13607	115.7961	-45.06803	1 vs 2	27.8851	<.0001

# Hypothesis testing 2. Random terms

Compare models with different random effects by a likelihood ratio test

```
> model11 <- lme(WEIGHT~SEX*TREAT, random=~1|BOX, data=gg)
> model0 <- lm(WEIGHT~SEX*TREAT, data=gg)
> anova(model11,model0)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
model11	1	8	89.34314	106.8604	-36.67157			
model0	2	7	119.63090	134.9585	-52.81545	1 vs 2	32.28776	<.0001

The model with the random component provides a significantly better fit.

# Treat box as random or fixed?

```
> model1 <- lme(WEIGHT~SEX*TREAT, random= ~1|BOX,  
  data=cool, method="ML")  
> model2 <- lm(WEIGHT~SEX*TREAT+BOX, data=cool)
```

```
> anova(model1,model2)
```

	Model	df	AIC	BIC	logLik
model1	1	8	78.25097	96.4643	-31.12548
model2	2	8	102.66355	120.8769	-43.33177

# Practical

- Day 4\_mixedeffects