
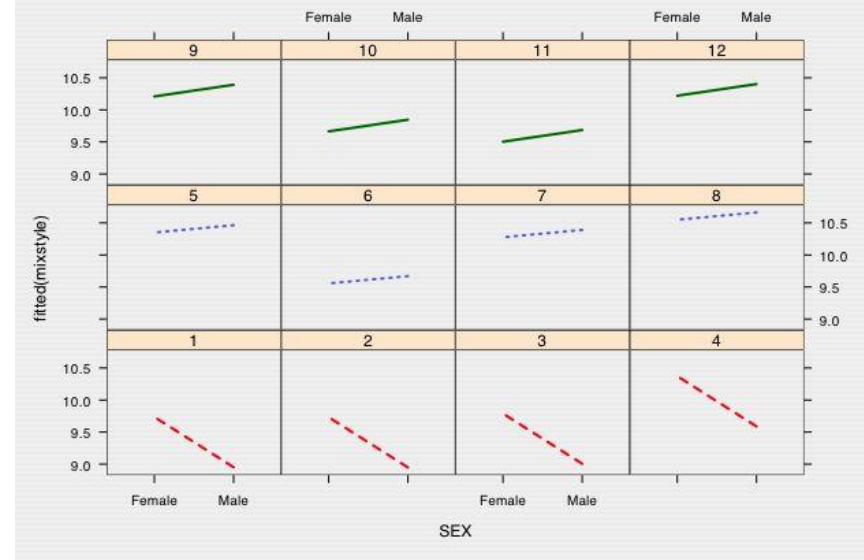


L8: More Random Models...



This morning...



- Intercepts were allowed to vary among nest boxes... under the influence of a treatment factor
- Model: $\text{weight} = a.\text{sex}_i + b.\text{temp}_j + \text{box}_j + \text{error}_i$
- Could the effect of sex vary among nest boxes as well?


Data = Mean + Treatment + Error + Group Treatment +
Group error + Group Treatment error

Example 2: Children's growth rates

A longitudinal dataset:

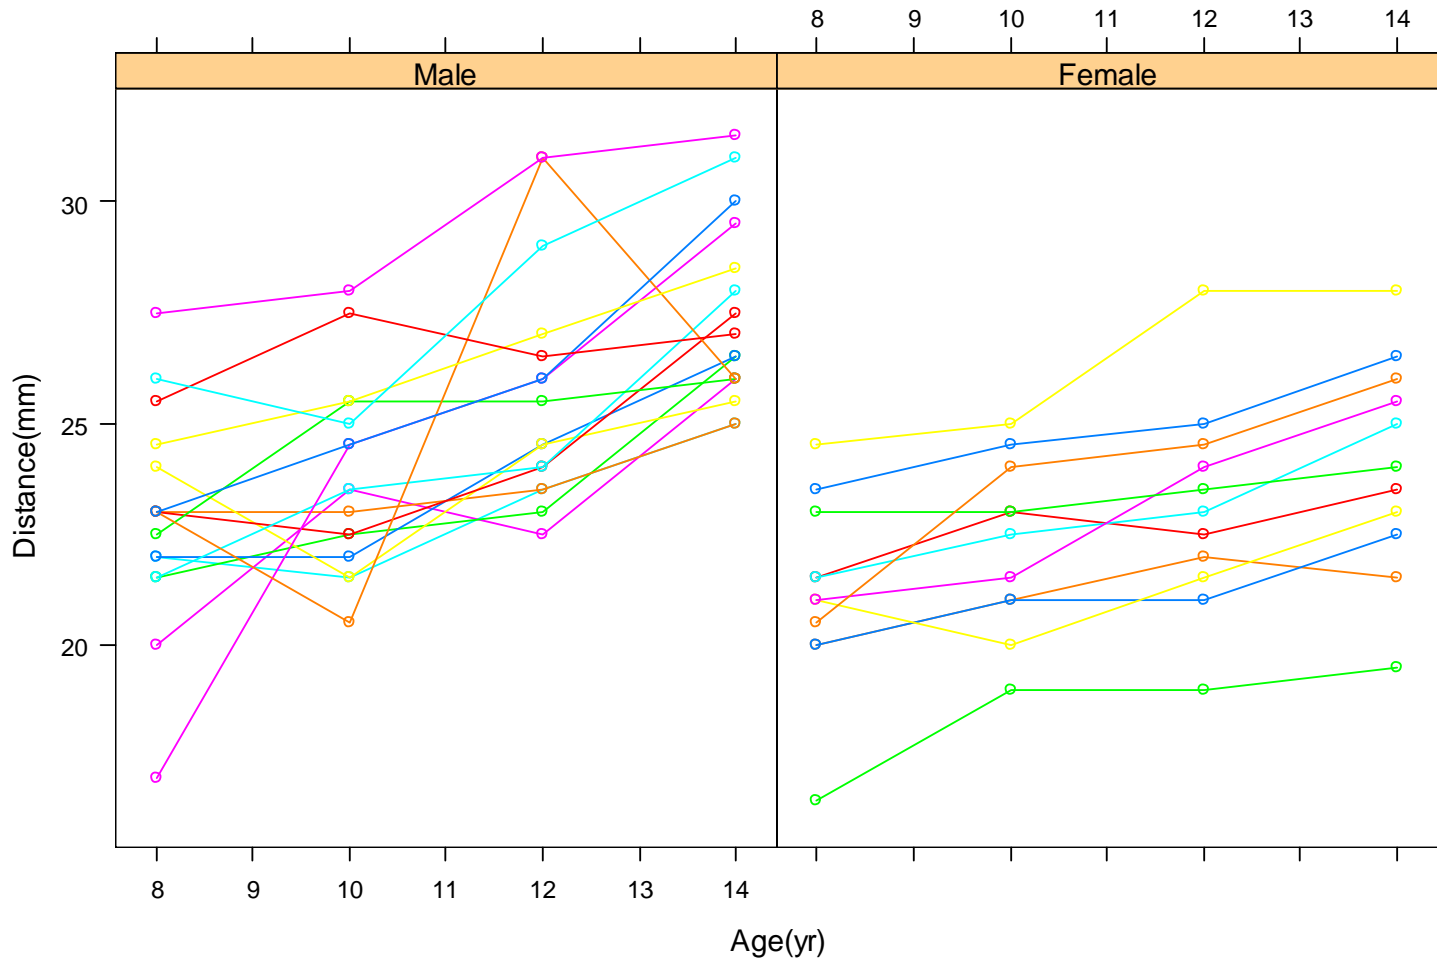
Data of skull measurements from children taken every two years from 8 until 14 from 27 children (16 males and 11 females).

- Orthodont dataset provided with library(nlme)
- Do the growth rates differ among the boxes?



`plot(Orthodont)`

Are **slopes** the same among boxes, within each sex?



Allow varying intercepts only:

```
> lme1 <- lme(distance~Sex*age, data=Orthodont, random= ~1|Subject)
```

Linear mixed-effects model fit by REML

Data: Orthodont

	AIC	BIC	logLik
	445.7572	461.6236	-216.8786

Random effects:

Formula: ~1 | Subject

(Intercept) Residual

StdDev: 1.816214 1.386382

Fixed effects: distance ~ Sex * age

	Value	Std.Error	DF	t-value	p-value
(Intercept)	16.340625	0.9813122	79	16.651810	0.0000
SexFemale	1.032102	1.5374208	25	0.671321	0.5082
age	0.784375	0.0775011	79	10.120824	0.0000
SexFemale:age	-0.304830	0.1214209	79	-2.510520	0.0141

:

:

Standardized Within-Group Residuals:

Min	Q1	Med	Q3	Max
-3.59804402	-0.45461690	0.01578364	0.50244657	3.68620794

Number of Observations: 108


Number of Groups: 27

Predictions with varying intercepts

```
plot( augPred(lme1, level=c(0:1)) )
```

The `level=` argument plots the predictions of

- the fixed effects for each sex (blue line) and
 - the fixed effect plus random effect for each individual (purple line)
- *Note that the random effect has allowed the **intercept** of the slope to differ for each individual from the predicted fixed effect, but **slope** for fixed and random effect is the same for each subject*



Allowing varying intercepts and slopes

```
> lme2 <- lme(distance ~ Sex*age, data=Orthodont, random=~age|Subject)
> summary(lme2)
```

Linear mixed-effects model fit by REML

Data: Orthodont

	AIC	BIC	logLik
	448.5817	469.7368	-216.2908

Random effects:

Formula: ~age | Subject

Structure: General positive-definite, Log-Cholesky

parametrization

	StdDev	Corr
(Intercept)	2.4059143	(Intr)
age	0.1804303	-0.668
Residual	1.3100046	

Fixed effects: distance ~ Sex * age

	Value	Std.Error	DF	t-value	p-value
(Intercept)	16.340625	1.0185753	79	16.042629	0.0000
SexFemale	1.032102	1.5958008	25	0.646761	0.5237
age	0.784375	0.0860090	79	9.119688	0.0000
SexFemale:age	-0.304830	0.1347502	79	-2.262183	0.0264

:

:

Standardized Within-Group Residuals:


	Min	Q1	Med	Q3	Max
	-3.167729145	-0.385826727	0.007027997	0.445020334	3.849531112

Number of Observations: 108

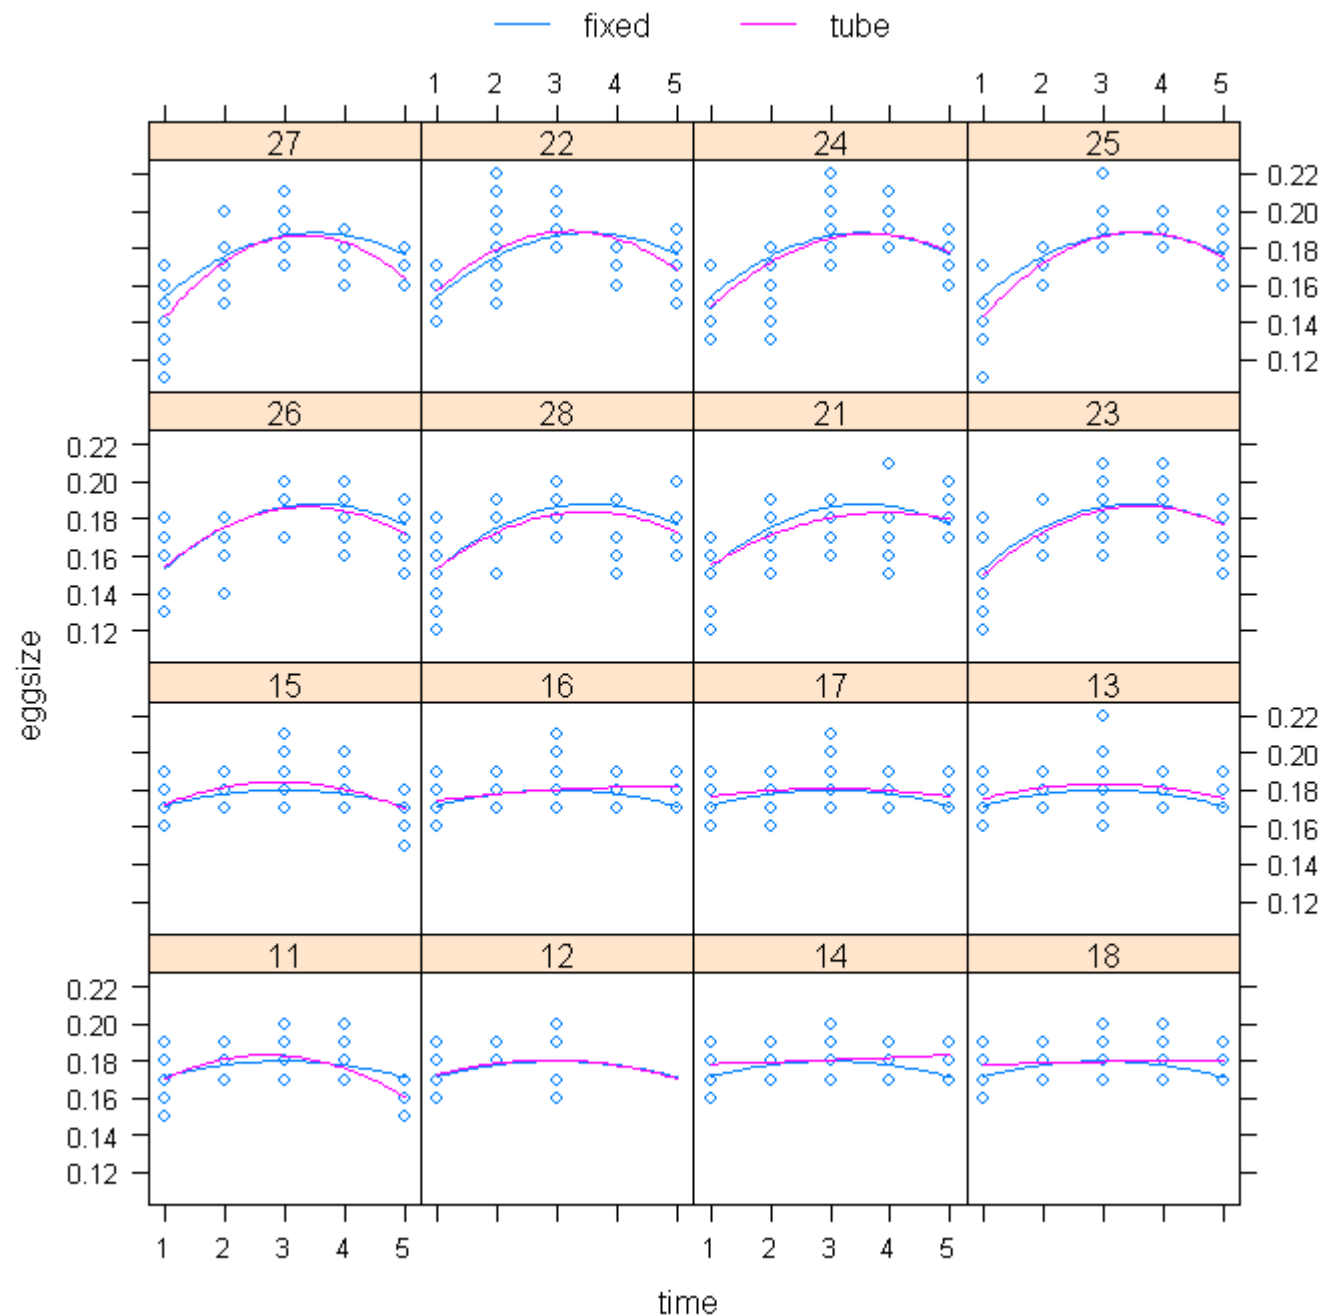
Number of Groups: 27

Predictions with varying intercepts and slopes

- Predictions for each subject now vary in both intercept and slope (purple line; compare blue line)



Paternal effects in mites



lme(
eggsize ~
poly(time,2) + food,

random =
~ maleID + poly(time,2)
| tube)

Hypothesis testing

- Do groups need both slopes and intercepts to vary?

```
lme1 <- lme( distance~Sex*age, data=Orthodont, random= ~age|Subject )  
lme2 <- lme( distance~Sex*age, data=Orthodont, random= ~1|Subject )
```

```
anova(lme1,lme2)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
lme1	1	8	448.5817	469.7368	-216.2908			
lme2	2	6	445.7572	461.6236	-216.8786	1 vs 2	1.175587	0.5556

→ No, a model with random intercepts appears sufficient.

- Is there a significant effect of sex on growth (i.e. the slope of the distance~age effect)?

```
anova(lme2)
```

	numDF	denDF	F-value	p-value
(Intercept)	1	79	4123.156	<.0001
Sex	1	25	9.292	0.0054
age	1	79	122.450	<.0001
Sex:age	1	79	6.303	0.0141

→ Yes, growth is significantly affected by sex.

Model structures
for mean and variance changing
in time and space


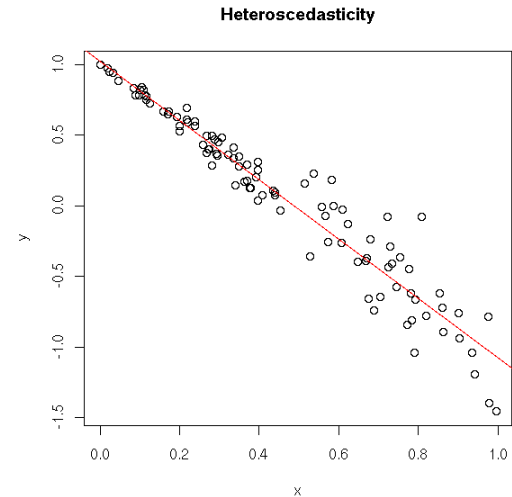
Complex model structures with library “nlme”

A “random factor” is a grouping variable.

- You may need to fit different **variances** for different groups.

“Random effects” may also include:


- **autocorrelation** if nearby replicates are more similar.



How do we model variance structures within a “regression” framework to incorporate changes in variance and autocorrelation?

(1) Variance changing with covariates

- Transform or fit appropriate model
- GLM with `quasi(link=ident,var="mu")`
- GLS with variance family
 - `gls()` if you don't have groups
 - Allows errors to be correlated and/or with different variances



varClasses

- **varIdent**: constant variance(s), generally used to allow different variances according to the levels of a classification factor.
- **varExp**: exponential of a variance covariate.
- **varPower**: power of a variance covariate.
- **varConstPower**: constant plus power of a variance covariate.
- **varFixed**: fixed weights, determined by a variance covariate.
- **varComb**: combination of variance functions.
- Eg:

```
var <- varPower(form= ~size)  
Mod1 <- gls(weight ~size, data=data, weights= var)
```


(2) Hierarchical grouping – spatial sampling designs

- E.g. repeated samples from different sites so you are looking at space as discrete

```
lme(plant_height ~ pesticide, random= ~1|farm/field)
```

Random effects structures:

(3) Correlation in time



- Temporal autocorrelation (TSA for another time)


```
> lme1 <- lme(distance ~ Sex*age, data=Orthodont, random=~1|Subject)
> lme3 <- update(lme1, correlation= corAR1())
> anova(lme3,lme1)
```

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
lme3	1	7	447.7081	466.2188	-216.8541			
lme1	2	6	445.7572	461.6236	-216.8786	1 vs 2	0.04913688	0.8246

This tests for a correlation in the resids between consecutive points (i.e. autocorrelation of lag 1)


(4) Correlation in space

```
wheat <- lme( yield~Treatment,  
  random= ~1|variety, data=Wheat3,  
  correlation= corSpher(c(28,0.2),  
    form= ~ latitude + longitude,  
    nugget= TRUE) )
```



corSpher(c(28,0.2)):

- 28 refers to range of correlations and
- 0.2 to the size of the *nugget* effect
 - correlation due to measurement error – how correlated are plants when close together? They should be about the same size, so plants measured at close points in space should be correlated at about zero, but it may be a small amount



corClasses

- **corAR1**: autoregressive process of order 1.
- **corARMA**: autoregressive moving average process, with arbitrary orders for the autoregressive and moving average components.
- **corCAR1**: continuous autoregressive process (AR(1) process for a continuous time covariate).
- **corCompSymm**: compound symmetry structure corresponding to a constant correlation.
- **corExp**: exponential spatial correlation.
- **corGaus**: Gaussian spatial correlation.
- **corLin**: linear spatial correlation.
- **corRatio**: Rational quadratics spatial correlation.
- **corSpher**: spherical spatial correlation.
- **corSymm**: general correlation matrix, with no additional structure.

Non-Normal Errors

Modelling non-normal responses

- GLMs are a generalisation of LMs where error structure is non normal
- What do you do with LMEs where errors are not normal?
 - GLMMs
 - `glmm.pql()` `library(MASS)` – penalised QuasiL
 - `lmer()` `library(lme4)`
 - Transform and use lme()?
- Active research field... To avoid believing silly answers from GLMM, do transformed LME first as guide

Table I. Techniques for GLMM parameter estimation, their advantages and disadvantages and the software packages that implement them

Technique	Advantages	Disadvantages	Software
Penalized quasilikelihood	Flexible, widely implemented	Likelihood inference inappropriate; biased for large variance or small means	PROC GLIMMIX (SAS), GLMM (Genstat), glmmPQL (R), glmer (R)
Laplace approximation	More accurate than PQL	Slower and less flexible than PQL	PROC GLIMMIX [56], glmer (R), glmm.admb (R), AD Model Builder, HLM
Gauss-Hermite quadrature	More accurate than Laplace	Slower than Laplace; limited to 2–3 random effects	PROC GLIMMIX [56], PROC NLMIXED (SAS), glmer (R), glmmML (R)
Markov chain Monte Carlo	Highly flexible, arbitrary number of random effects; accurate	Very slow, technically challenging, Bayesian framework	WinBUGS, JAGS, MCMCpack, (R), AD Model Builder

Review



Generalized linear mixed models: a practical guide for ecology and evolution

Benjamin M. Bolker¹, Mollie E. Brooks¹, Connie J. Clark¹, Shane W. Geange², John R. Poulsen¹, M. Henry H. Stevens³ and Jada-Simone S. White¹

¹ Department of Botany and Zoology, University of Florida, PO Box 118525, Gainesville, FL 32611-8525, USA

² School of Biological Sciences, Victoria University of Wellington, PO Box 600, Wellington 6140, New Zealand

³ Department of Botany, Miami University, Oxford, OH 45056, USA

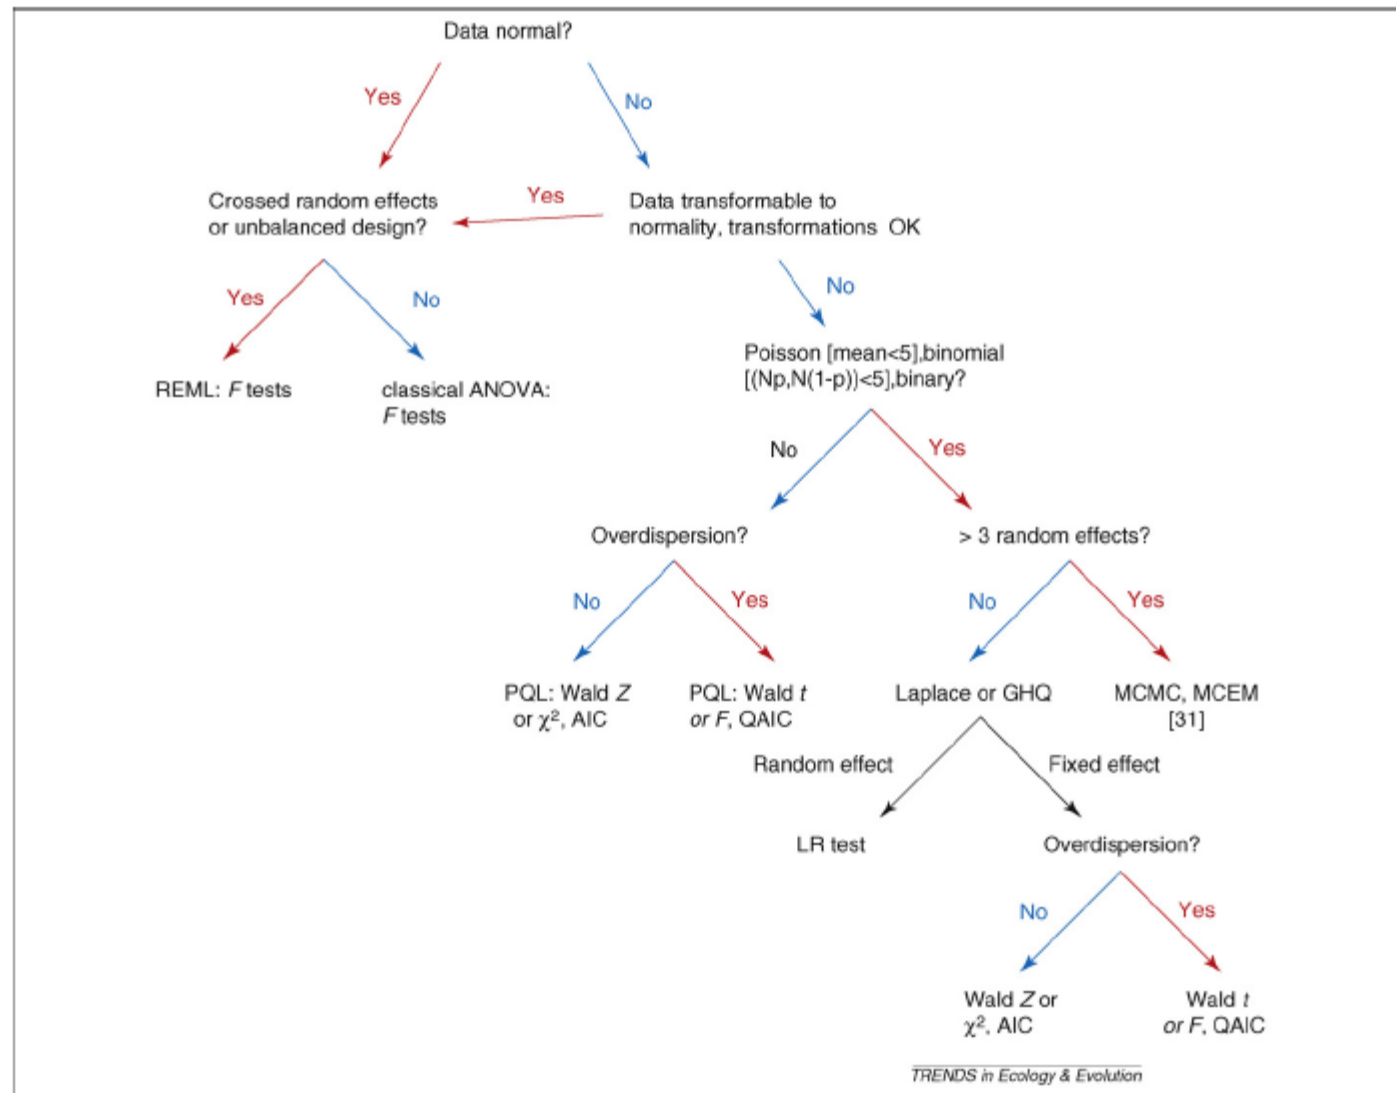


Figure 1. Decision tree for GLMM fitting and inference. Conditions on the Poisson and binomial distributions along the right branch refer to penalized quasikelihood (PQL) rules of thumb [30]: to use PQL, Poisson distributions should have mean > 5 and binomial distributions should have the minimum of the number of successes and failures > 5 . MCEM = Monte Carlo expectation-maximization [40].

Modelling non-normal responses

Various options in R (still a research topic)

```
> library(lme4)
```

```
> lmer1 <- lmer(y ~ x + (1|z), family=
  binomial)
```

- where z is the random factor

```
> library(MASS)
```

```
> glmm1 <- glmmPQL(distance ~ Sex*age,
  data=Orthodont, random= ~1|Subject,
  family="gaussian")
```

Linear mixed-effects model fit by REML

Data: Orthodont
AIC BIC logLik
445.7572 461.6236 -216.8786

Random effects:
Formula: ~1 | Subject
(Intercept) Residual
StdDev: 1.816214 1.386382

Fixed effects: distance ~ Sex * age

	Value	Std.Error	DF	t-value	p-value
(Intercept)	16.340625	0.9813122	79	16.651810	0.0000
SexFemale	1.032102	1.5374208	25	0.671321	0.5082
age	0.784375	0.0775011	79	10.120823	0.0000
SexFemale:age	-0.304830	0.1214209	79	-2.510520	0.0141

```
> glmm1 <- glmmPQL(  
  distance~Sex*age,data=  
  Orthodont,  
  random=~1|Subject,  
  family="gaussian")  
> summary(glmm1)
```

```
> lmer1 <- lmer(distance~sex*age  
  +(1|Subject),  
  family="gaussian",  
  data=Orthodont)  
> summary(lmer1)
```

Linear mixed-effects model fit by maximum likelihood
Data: Orthodont
AIC BIC logLik
NA NA NA

Random effects:
Formula: ~1 | Subject
(Intercept) Residual
StdDev: 1.740851 1.369159

Variance function:
Structure: fixed weights
Formula: ~invwt
Fixed effects: distance ~ Sex * age

	Value	Std.Error	DF	t-value	p-value
(Intercept)	16.340625	0.9814310	79	16.649795	0.0000
SexFemale	1.032102	1.5376069	25	0.671239	0.5082
age	0.784375	0.0779963	79	10.056564	0.0000
SexFemale:age	-0.304830	0.1221968	79	-2.494580	0.0147

Generalized linear mixed models: a practical guide for ecology and evolution

Benjamin M. Bolker¹, Mollie E. Brooks¹, Connie J. Clark¹, Shane W. Geange², John R. Poulsen¹, M. Henry H. Stevens³ and Jada-Simone S. White¹

¹ Department of Botany and Zoology, University of Florida, PO Box 118525, Gainesville, FL 32611-8525, USA

² School of Biological Sciences, Victoria University of Wellington, PO Box 600, Wellington 6140, New Zealand

³ Department of Botany, Miami University, Oxford, OH 45056, USA

Trends in Ecology and Evolution Vol.24 No.3


127-135 (2008)

Alain F. Zuur • Elena N. Ieno
Neil J. Walker • Anatoly A. Saveliev
Graham M. Smith

Mixed Effects Models and Extensions in Ecology with R

Zero Inflated Models and Generalized Linear Mixed Models with R

Alain F. Zuur
Anatoly A. Saveliev
Elena N. Ieno



Data Analysis Using Regression and Multilevel/Hierarchical Models

ANDREW GELMAN
JENNIFER HILL